

Creating Singularity containers for HPC users

Martin Čuma
Center for High Performance Computing
University of Utah
m.cuma@utah.edu

- Why do we want to use containers?
- Containers basics
- Prepare your computer for containers
 - As a backup ssh to our lab
- Build and deploy a container
- Containers for complex software
- Containers for Windows programs

1. Download the talk slides

http://home.chpc.utah.edu/~mcuma/chpc/Containers_RMACC17.pdf

2. Get an user/password paper slip

3. Using terminal application (Mac terminal, PuTTY, GIT Shell)

- `ssh userxx@linuxclass.chpc.utah.edu`


4. Make sure you can see singularity

- `ls /usr/local/bin/singularity`

5. Make sure you can sudo singularity command

- `sudo /usr/local/bin/singularity -version`

OR – if you have installed Singularity on your laptop, use it




Why to use containers?

- Some programs require complex software environments
 - OS type and versions
 - Drivers
 - Compiler type and versions
 - Software dependencies
 - Python/R/MATLAB versions
 - glibc, stdlibc++ versions
 - Other libraries and executables
 - Python/R libraries
- We may not want to build everything from a scratch

- Research outputs include software and data
- Software reproducibility
 - Software repositories (svn, git)
 - Good but often software has dependencies
- Data reproducibility
 - Data as publication supplementary info, centralized repositories (NCBI), ...
 - Disconnected from the production environment
- Package data AND code AND compute environment in one file

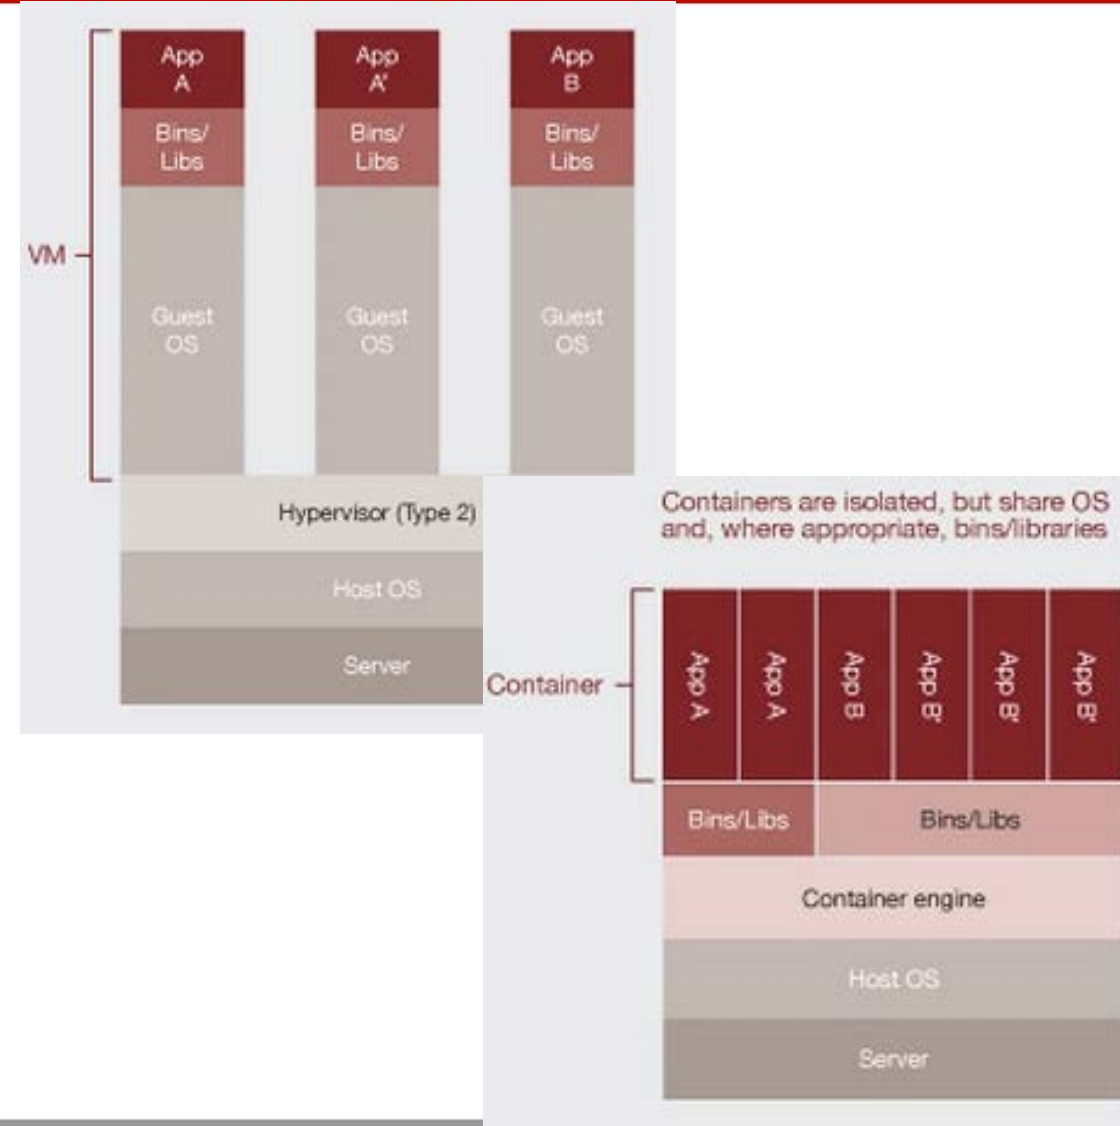
- Develop a program / pipeline locally, run globally
- Scale to parallel resources
 - Run many times
 - Use local or national HPC resources
- Automate the process
 - Container/software building and deployment
 - Parallel pipeline

- Old applications on old Linux versions



Container basics

- Hardware virtualization
 - Running multiple OSes on the same hardware
 - VMWare, VirtualBox
- OS level virtualization
 - run multiple isolated OS instances (guests) under a server OS (host)
 - Also called containers
 - Docker, Singularity



- Isolate computing environments
 - And allow for regenerating computing environments
- Guest OS running over host OS
 - Guest's OS can be different than host's
 - Low level operations (kernel, network, I/O) run through the host
- From user standpoint guest OS behaves like standard OS

- Docker
 - Well established
 - Has docker hub for container sharing
 - Problematic with HPC
- Singularity
 - Designed for HPC
 - Support for SLURM, MPI, GPUs
 - Relatively new but rapidly expanding
- Shifter, Charliecloud
 - HPC oriented but less used presently





- Integrate with traditional HPC
 - Same user inside and outside of the container
 - Same file systems (home, scratch), environment
 - Can integrate with existing software (CHPC sys branch)
- Portable and sharable
 - A container is a file
 - It can be built on one OS and run on another
- Only Linux support right now
 - But any Linux version – RHEL/CentOS, Ubuntu, Debian

Container Creation

```
sudo singularity create container.img
```

Import and Bootstrap

```
sudo singularity import container.img docker://ubuntu
```

```
sudo singularity bootstrap container.img ubuntu.def
```

Interact and Modify

```
sudo singularity shell --writable container.img
```

USER ENDPOINT

Container Execution

```
singularity run container.img
```


```
singularity shell container.img
```

```
singularity exec container.img ...
```

SHARED COMPUTATIONAL RESOURCE

./container.img

Kurtzer GM, Sochat V, Bauer MW. Singularity: Scientific Containers for Mobility of Compute (under review)



Prepare your computer for Singularity containers


- We need to run Linux to build/run Singularity
 - If you already run Linux, make sure you have a root and then take a nap
 - On Windows and Mac, we need to install Linux first
- Install Linux in a VM
 - Windows – GIT Bash, Virtual Box and Vagrant
 - <https://www.chpc.utah.edu/documentation/software/containers-localbuild.php#prep>
 - Mac – Homebrew with Virtual Box and Vagrant
 - <http://singularity.lbl.gov/install-mac>

- Windows – GIT Bash, VirtualBox, Vagrant
 - GIT Bash provides a bash terminal on Windows
 - VirtualBox provides VM virtualization
 - Vagrant automates VM setup
 - Use Vagrant to install Ubuntu VM
- Mac – VirtualBox and Vagrant
 - Already have a terminal
 - Use Homebrew to install VirtualBox and Vagrant

<https://www.chpc.utah.edu/documentation/software/containers-localbuild.php#prep>

- Linux laptop users – wake up
- In Ubuntu VM, or standalone Linux

```
$ git clone https://github.com/singularityware/singularity.git
$ cd singularity
$ ./autogen.sh
$ ./configure --prefix=/usr/local
$ make
$ sudo make install
```



Build and run containers

- Building a container requires a root, or sudo
 - You can do that on your own machine
 - You can't do that at CHPC clusters
 - > build your containers locally
 - (we make an exception on our linuxclass machine)
- You can run a container as an user
 - You can run your own containers at CHPC
 - You can run CHPC provided containers at CHPC



- Singularity allows to run images from Docker hub (and Singularity hub)

```
$ singularity shell docker://ubuntu:latest
```

```
$ whoami
```

```
$ env | grep SINGULARITY
```

```
$ exit
```

- Other ways to run

```
$ singularity exec image program
```

```
$ singularity run image
```

- All needs to be run as root or with sudo
- Create a bare container

```
$ sudo /usr/local/bin/singularity create --size 2048
```

- Install (bootstrap) the container

```
$ sudo /usr/local/bin/singularity bootstrap ubuntu16.img ubuntu16.def
```

- If additional installation is needed after the bootstrap

- Shell into the container and do the install manually

```
$ sudo /usr/local/bin/singularity shell -w -s /bin/bash ubuntu16.img
```

- Execute a script in the container

```
$ sudo /usr/local/bin/singularity exec -w ubuntu16.img myscript.sh
```

- Defines how the container is bootstrapped
 - Header – defines the core OS to bootstrap
 - Sections – scriptlets that perform additional tasks
- Header
 - Docker based (faster installation)

```
BootStrap: docker
```

```
From: ubuntu:16.10
```

- Linux distro based

```
BootStrap: debootstrap
```

```
OSVersion: xenial
```

```
MirrorURL: http://us.archive.ubuntu.com/ubuntu/
```

- `%setup` Runs on the host
 - Install host based drivers (e.g. GPU)
- `%post` Runs in the container
 - Install additional packages, configure, etc
- `%runscript` Defines what happens when container is run
 - Execution commands
- `%test` Runs tests after the bootstrap
 - Basic testing

- Download CHPC containers GIT repo

```
$ git clone https://github.com/CHPC-UofU/Singularity-ubuntu-python
```

Go to the ubuntu_python directory and view what's in there

```
$ cd https://github.com/CHPC-UofU/Singularity-ubuntu-python
$ ls
$ cat build_container.sh
... replace singularity with /usr/local/bin/singularity
$ more Singularity
```

- Simply type the build script

```
$ ./build_container.sh
```

- CHPC specific caveats
 - File server mount points

```
$ mkdir /uufs /scratch
```

- Locally

```
$ singularity shell ubuntu_python.img
```

```
$ /usr/bin/python -c "import numpy as np;np.__config__.show( )"
```

- At our linuxclass machine:

```
$ scp ubuntu_python.img userxx@linuxclass.chpc.utah.edu:~/
```

```
$ ssh userxx@linuxclass.chpc.utah.edu
```

```
$ singularity shell ubuntu_python.img
```

```
Singularity $ /usr/bin/python -c "import numpy as  
np;np.__config__.show( )"
```

- Binding mount points


```
$ export SINGULARITY_BINDPATH="/scratch,/uufs/chpc.utah.edu"  
$ singularity shell -B /scratch,/uufs/chpc.utah.edu  
ubuntu_python.img
```

- Specifying shell

```
$ export SINGULARITY_SHELL=/bin/bash  
$ singularity shell -s /bin/bash ubuntu_python.img
```

- More specialized topics – ask us

- Using environment modules from the host
- Using GPUs, MPI over InfiniBand



- Many Linux programs are binary compatible between distros
 - Most installed binaries are (Intel, PGI tools, DDT, ...)
- No need to install these in the container – use our NFS mounted software stack through Lmod
 - Need to have separate Lmod installation for Ubuntu due to some files having different location
- In the container
 - Install Lmod dependencies
 - Modify /etc/bash.bashrc to source our Lmod


<https://github.com/CHPC-UofU/Singularity-ubuntu-python/blob/master/Singularity>

- Need to bring in the Nvidia driver stack
 - Pre Singularity 2.3 – explicitly install – make sure to have the same driver version on the host and in the container
 - Singularity 2.3+ - -nv flag

<https://github.com/CHPC-UofU/Singularity-tensorflow/blob/master/Singularity>


- Need to bring in the InfiniBand driver stack
 - For Ubuntu based on <https://community.mellanox.com/docs/DOC-2431>

<https://github.com/CHPC-UofU/Singularity-tensorflow/blob/master/Singularity>




Containers for complex software

- Complex software dependencies
 - Especially Python and R packages
 - bioBakery – intricate dependencies of Python and R which did not build on CentOS
 - SEQLinkage – instructions to build on Ubuntu using its packages
- Quick deployment
 - Some Linux distros provide program packages while others don't
 - paraview-python on Ubuntu via apt-get
- Deploying your own code or pipeline




- Bootstrap the basic container
- Shell into the container
 - Install additional needed programs
 - If they have dependencies, install the dependencies – google for the OS provided packages first and install with apt-get/yum if possible
 - Put the commands in the `%post` scriptlet
- Build the container again
 - Now with the additional commands in the `%post`
 - If something fails, fix it, build container again
- Iterate until all needed programs are installed

- Instructions say to
 - Install VirtualBox, Vagrant, and bioBakery from an archive
 - Great for a desktop, but, not for an HPC cluster
 - Further below they mention Google Cloud
- So we download the bioBakery archive, unpack it and look inside
 - Great, there is `google_cloud/build_biobakery.sh` script
 - In that file, Ubuntu 16.04 is mentioned




- Build base Ubuntu 16.04 container
 - sudo shell into the container
 - Start executing the lines of the build_biobakery.sh script, one after another
 - Some dependencies pop up, install them
 - Another caveat – Linuxbrew requires to be installed as non-root
 - Do some web searching and figure how to add a new user and run Linuxbrew as this user
 - In the end, add the correct paths to the container environment
- ```
$ echo "export PATH=/usr/local/bin:$PATH" >> /environment
```



- Once everything installs in the container
  - Run the bioBakery tests
  - Add %test section that run the bioBakery tests
  - Build the container again, now it will run the tests (will take a few hours)
- Create a module file or an alias to start the container
- See it all at

<https://github.com/CHPC-UofU/Singularity-bioBakery>


- <http://singularity.lbl.gov>
- <https://singularity-hub.org>
- [https://www.chpc.utah.edu/documentation/software/container\\_s.php](https://www.chpc.utah.edu/documentation/software/container_s.php)
- [https://github.com/mcuma/chpc\\_singularity](https://github.com/mcuma/chpc_singularity)



# Windows in a container?

- What, Windows?
  - There are programs that researchers use that only run on Windows
  - E.g. data processing that comes with an instrument
- Our current approach
  - Tell them to run on our only Windows server
    - Gets oversubscribed quickly
  - Build a specific VM
    - Resource intensive for us, not high performing
- What if we could run Windows programs on our Linux clusters





- Windows compatibility layer on Linux
  - <https://www.winehq.org/>
  - Not an emulator – translates Windows system calls to Linux, provides alternative Windows system libraries,...
  - Actively developed, under CodeWeavers company umbrella
  - Windows ABI completely in user space
  - Most Linux distros come with some version of Wine
  - Generally better to use recent Linux distros for more recent Wine version (<https://www.winehq.org/download>)

- While Wine provides the basic Windows support, Winetricks is a set of scripts that install additional Windows libraries
  - Like library dependencies in Linux
  - `wintricks list` – to list available libraries
  - Most commonly used libraries are DirectX, .NET, VB or C runtimes

- Poached out of <http://dolmades.org/>
- Basic Singularity container
  - Recent Ubuntu or Fedora
  - Some winetricks work better on Fedora than Ubuntu, and vice versa
  - Include Wine repo from winehq to get the latest Wine version
  - Some experimentation is needed but if the Windows program is not complicated, success chances are there

- Install Wine and Winetricks

```
dpkg --add-architecture i386
```

```
apt update
```

```
apt -y install wget less vim software-properties-common
python3-software-properties apt-transport-https winbind
```

```
wget https://dl.winehq.org/wine-builds/Release.key
```

```
apt-key add Release.key
```

```
apt-add-repository https://dl.winehq.org/wine-builds/ubuntu/
```

```
apt update
```

```
apt install -y winehq-stable winetricks
```

- User application
  - Done in `%runscript` section
  - First container launch creates WINEPREFIX (Windows file space), then installs the needed applications, and tars the whole WINEPREFIX for future use
  - Subsequent container launch untars WINEPREFIX and launches program

```
TEMPDIR="$(mktemp -d)"
APPDIR="$HOME/WINE/Topofusion"
PROFILEDIR="$HOME/WINE/PROFILES/${USER}@${HOSTNAME}"
...
export WINEPREFIX="$TEMPDIR/wineprefix"
export WINEARCH="win32"


if [-f "$APPDIR/wineprefix.tgz"]; then
 echo "Found existing wineprefix - restoring it..."
 mkdir -p "$WINEPREFIX"; cd "$WINEPREFIX"; tar xzf "$APPDIR/wineprefix.tgz"
else
 wineboot --init
 echo "Installing TopoFusion and its dependencies ..."
 winetricks dlls directx9 vb6run
 wget http://topofusion.com/TopoFusion-Demo-Pro-5.43.exe
fi
wine ./TopoFusion-Demo-Pro-5.43.exe
```

- IDL 6.4 runtime + PeakSelector
  - IDL runtime under Linux crashes due to IDL bug
  - Windows runtime works fine, older IDL (ca. 2010)
  - <https://github.com/CHPC-UofU/Singularity-ubuntu-wine-peakselector>
- Topofusion
  - My favorite GPS mapping program, e.g.  
<http://home.chpc.utah.edu/~mcuma/summer16/madison/wed/>
  - Needs DirectX and VB runtime
  - <https://github.com/CHPC-UofU/Singularity-ubuntu-wine-topofusion>

- Very new application (Win10 like)
  - Installer was not functional under Wine
- Complex scientific application
  - .NET – did not install on Ubuntu, worked on Fedora
  - Microsoft SQL did not install – show stopper
- Wine application compatibility
  - <https://appdb.winehq.org/>
  - Notice a lot of games



- Success rate 1 out of 3 is not that great
  - Still worth trying, the chances are there
  - Singularity makes it easier to experiment
- It would be nice to have a HPC support for Windows so that
  - We would not need to have specialized Win machines
  - We would not have to build special purpose VMs
- May still need to look into the direction of reconfigurable HPC clusters like Bridges or Jetstream



# Open discussion

- What do you do with complicated installs?
- What is your experience with containers?
- How are you supporting other Oses?