

How to Access Summit

Shelley Knuth

shelley.knuth@colorado.edu

www.rc.colorado.edu

Slides: <https://github.com/ResearchComputing/RMACC/2017>

Outline

- What is RC?
- Steps to get access to our systems
 - Accounts
 - Two-factor authentication
 - Allocations
 - Logging in
- Navigating our systems
 - Storage spaces
 - Data transfer - Globus
 - Software
- Running jobs

Hardware - Summit Supercomputer

- 475 compute nodes (Intel Xeon Haswell)
 - 24 cores per node
 - 11,400 total cores
 - Omni-Path network
 - 1.2 PB scratch storage
 - GPFS File system
-
- 67% CU, 23% CSU, 10% RMACC



Additional Types of Summit Compute Nodes

- 10 Graphics Processing Unit (GPU) Nodes
 - NVIDIA Tesla K80 (2/node)
- 5 High Memory Nodes
 - 2 TB of memory/node, 48 cores/node
- Phi Nodes
 - 20 nodes
 - Intel Xeon Phi

How To Access RC Resources?

1. Get an account
 2. Set up two factor authentication
 3. Set up an allocation
 4. Log in
 5. Create greatness
- After you login, you will need to do many additional things that we will discuss today

Getting an RC Account

- CU Boulder users and affiliates:
 - Request an account through the RC Account request portal
 - <https://portals.rc.colorado.edu/accounts/account-request/create>
- CSU Users:
 - Request an CSU eID if you don't have one
 - Fill out account application form
 - Duo authentication
 - Then get an RC user account
 - <https://www.acns.colostate.edu/hpc/summit-get-started/>
- RMACC Users:
 - We are working on a login system through XSEDE
 - Not operational yet (as of this writing)

Setting up Two-Factor Authentication

- Two factor authentication is required to access our system
- Require this to provide an extra level of authentication
- Two methods for achieving this:
 - Duo
 - Access through a smart phone app
 - Vasco OTP (one time password)
 - CU only
 - Physical device

Duo Authentication

- Once you get an account, contact rc-help@colorado.edu to request a Duo invitation
- Once you get the invitation, you'll get a series of steps to complete Duo enrollment
- RC supports Duo “push” and “phone call” for authentication
- Greatly prefer “push”

Vasco Authentication

- Physical device that generates a new password every 30 seconds
- To get the device, go to the IT Service Center or request a time to pick one up from Research Computing
- You will need to show an ID to get a device
- First, register the device
 - Otp.colorado.edu
- You will set a four digit pin

Allocations

- You will need a compute allocation to use any of our resources
- Currently, to request an allocation please email rc-help@colorado.edu and ask for a General allocation
 - Need to provide a few sentences on your project
- In the future, we will have a place on our website to submit a more formal request
- Once you have some benchmarks, you will want to move to a project allocation
- Track use on system
 - Reporting
 - Enough resources

What is Fair Share?

- Fair share scheduling uses a complex formula to determine priority in queue
- Looks at load for each user and each QOS and balances utilization to fairly share resources
 - Involves historical use by user plus how long job has been in the queue
- System will first look at weighted average utilization of user over last 4 weeks
- Then compare it to the fair share target percentage of a user

Fair Share Target Percentage

- The target percentage depends on your priority based on your project proposal
- Everyone not associated with a project at CU for example shares a target percentage of 13% (20% of the CU fraction)
 - No guaranteed level per user
- If you are under (over) your target percentage (based on a 4 week average) your priority is increased (decreased)
- Reminder this all only impacts pending jobs
- If no other pending jobs and enough resources are available then your job will run regardless of your previous usage

Transferring Data

- Globus is Research Computing's preferred method of data transfer
- Designed with researchers in mind
- End points between computers make for efficient data transfer with an easy to use interface
 - Endpoints are different locations that data can be moved to/from
 - Personal or multi-user
- Rsync and sftp through the login nodes is good for small transfers
- www.globus.org

Software

- Common software is available to everyone on the systems
- Can install your own software
 - But you are responsible for support
 - We are happy to assist
- Research Computing uses modules to manage software
 - You can load modules to prepare your environment for using software
 - Set any environment variables
 - Set environment so application can find appropriate libraries, etc.

Important Things to Know About Modules

- Some modules might require a specific hierarchy to load
 - For some modules, you may need to specify a specific version
 - For example, **module load R/3.3.0**
 - For other modules, you may be able to be more generic
 - For example, **module load matlab**
- Some modules may require you to first load other modules that they depend on
- To find prerequisites for a module, type `module spider <package>`
- To find out what software is available, you can type **module avail**
- To set up your environment to use a software package, type **module load <package>/<version>**

Job Submission

Running Jobs

- What is a “job”?
- Interactive jobs
 - Work interactively at the command line of a compute node
- Batch jobs
 - Submit job that will be executed when resources are available
 - Create a text file containing information about the job
 - Submit the job file to a queue
- We use SLURM to manage resources and schedule jobs

Partitions and ‘Quality of Services’

- There are several ways to define where your job will run
- Partitions (basically a queue):
 - Resources/hardware
- QoS:
 - Tells what the limits or characteristics of a job should be
 - Maximum wall time
 - Number of nodes
- One partition might have multiple QoS
- A QoS might exist on multiple partitions

Available Partitions

Partition	Description	# of nodes	cores/node	GPUs/node
shas	General Compute (Haswell)	380	24	0
sgpu	GPU-enabled nodes	10	24	effectively 4
smem	High-memory nodes	5	48	0
sknl	Phi (Knights Landing) nodes	20	68	0

Quality of Service

QoS	Description	Maxwall	Max jobs/user	Max nodes/user
normal	Default QoS	Derived from partition	n/a	256
debug	For quick turnaround when testing	1 H	1	32
long	For jobs needing longer wall times	7 D	n/a	20
condo	For groups who have contributed to the Summit condo	7 D	n/a	n/a

Useful Slurm Commands - sbatch

- **sbatch:** submit a batch script to slurm
- You can use a bunch of flag options in a batch script or on the command line
- Useful to put in script so have for future use

- Example:

```
sbatch test.sh
```

OR

```
sbatch --partition=shas test.sh
```

<http://slurm.schedmd.com/sbatch.html>

SBATCH Options

<http://slurm.schedmd.com/sbatch.html>

```
#SBATCH <options>          sbatch <options>
```

- Allocation: `--account=<account_no>`
 - Partition: `--partition=<partition_name>`
 - Sending emails: `--mail-type=<type>`
 - Email address: `--mail-user=<user>`
 - Number of nodes: `--nodes=<nodes>`
 - Number of tasks: `--ntasks=<processes>`
 - Quality of service: `--qos=<qos>`
 - Reservation: `--reservation=<name>`
 - Wall time: `--time=<wall time>`
 - Job Name: `--job-name=<jobname>`
- FYI: You do NOT actually type <> above – this designates something specific you as a user must enter about your job

Working on Summit

- Make sure you load the appropriate slurm module

```
module load slurm/summit
```

- After you run this command you can run sbatch to submit jobs

Practice Examples

Submit Your First Job!

- Submit a slurm job with the following instructions:
 1. The job should run the Unix “hostname” command
 2. The job will be submitted from a bash script named `hostname_summit.sh`
 3. The job will run on 1 node
 4. We will request 1 minute wall time
 5. Run from the debug QOS
 6. Run on the shas partition
 7. Use the rmacc reservation
 - This is only for this workshop

Hostname_summit.sh

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=0:01:00
#SBATCH --qos=debug
#SBATCH --partition=shas
#SBATCH --output=hostname_%j.out
#SBATCH --reservation=rma

# Number of requested nodes
# Max wall time
# Specify debug QOS
# Specify Summit haswell nodes
# Rename standard output file
# Reservation name

# purge all existing modules
module purge

hostname
```

Running the script

- Load up the slurm module
`module load slurm/summit`
- Submit the job:
`sbatch hostname_summit.sh`
- Check output

Another slurm command

- **squeue**
 - View information about jobs located in the slurm scheduling queue
- **OPTIONS:**
 - User: `-u <user_list>`
 - Queues: `--qos=<qos_list>`
- **EXAMPLE:**
`squeue --qos=debug`

<http://slurm.schedmd.com/squeue.html>

Your turn

- Submit a slurm job with the following instructions:
 1. The job should run first the whoami command, then the Unix “sleep” command for 30 seconds, then the hostname command
 - Syntax for these Unix commands are below:

```
whoami
```

```
sleep 30
```

```
hostname
```

<http://slurm.schedmd.com/squeue.html>

Your turn

- Submit a slurm job with the following instructions:
 1. The job will be submitted from a bash script named sleep.sh
 2. The job will run on 1 node
 3. Request a 1 minute wall time
 4. Run the job from the normal QOS
 5. Run the job from the Summit haswell partition
 6. Name your job sleep
 7. Email yourself the results at the end of the job run
 - Hint: Requires two SBATCH options to do this – see link at top of this slide
 8. Use the rmacc reservation

Sleep.sh

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=0:01:00
#SBATCH --qos=normal
#SBATCH --partition=shas
#SBATCH --output=sleep_%j.out
#SBATCH --job-name=sleep
#SBATCH --mail-type=end
###SBATCH --mail-user=<user>@colorado.edu
#SBATCH --reservation=rmacc

# Number of requested nodes
# Max walltime
# Specify normal QOS
# Specify Summit GPU nodes
# Rename standard output file
# Job submission name
# Email you when the job ends
# Email address to send to
# Reservation name

# purge all existing modules
module purge

whoami
sleep 30
hostname
```

Running an external script

- Let's run a Matlab program
- We will run the batch script `matlab.sh`
- This script calls and runs `matlab_tic.m`

Running the script

- Submit the job:

```
sbatch matlab.sh
```

- Check output

Matlab.sh

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=0:02:00
#SBATCH --qos=debug
#SBATCH --partition=shas
#SBATCH --output=matlab_%j.out
#SBATCH --reservation=rma

# purge all existing modules
module purge

# Load Matlab module
module load matlab

# Run matlab without a GUI
matlab -nodisplay -nodesktop -r "clear; matlab_tic;"
```

```
# Number of requested nodes
# Max walltime
# Specify debug QOS
# Specify Summit haswell nodes
# Output file name
# Reservation name
```

Your turn

- Submit a slurm job with the following instructions:
 1. Create an R program called `R_program.R` that creates a vector called “planets” and then list the planets in the vector
 - Syntax: `planets -> planets <- c("Mercury", "Venus", "Earth", "Mars", "Jupiter", "Saturn", "Uranus", "Neptune", "Pluto")`
 2. Print off the vector
 - Syntax: `planets`
 3. Create a bash script called `R_code.sh` that runs the R script
 - Syntax: `Rscript R_code.sh`
 4. The job will run on 1 node
 5. We will request a 1 minute wall time
 6. Specify the debug QOS
 7. Specify the shas partition
 8. The output will be put in a file called `R_code_%j.out`
 9. Use the rmacc reservation
 10. Don't forget to load the R module!

Solution – R_code.sh

```
#!/bin/bash
#SBATCH --nodes=1
#SBATCH --time=0:01:00
#SBATCH --qos=debug
#SBATCH --partition=shas
#SBATCH --output=R_code_%j.out
###SBATCH --reservation=rmacc

# Number of requested nodes
# Max walltime
# Specify debug QOS
# Specify Summit haswell nodes
# Output file name
# Reservation name

# purge all existing modules
module purge

# Load the R module
module load R/3.3.0

# Run R Script
Rscript R_program.R
```

Solution – R_program.R

```
#Simple R code example by Shelley Knuth (shelley.knuth@colorado.edu)

# Create vector
planets <- c("Mercury", "Venus", "Earth", "Mars", "Jupiter", "Saturn", "Uranus",
"Neptune", "Pluto")

# Print off vector
planets
```

Interactive jobs!

- Sometimes we want our job to run in the background
- Sometimes we want to work in program in real time
- For example, Matlab
- Let's run an interactive Matlab job

Interactive job

- To do this, we are going to log out and log back in
 - Only necessary for demo
 - Need to add something to the sign in process
- For Mac Users:
`ssh -X username@tutorial-login.rc.colorado.edu`
- For Windows Users, must set up X-forwarding through your SSH client program
- Also must have an X-server package on your laptop
 - Xming for Windows or XQuartz for Mac

Interactive job

- To work with Matlab interactively, we're going to request some time from the supercomputer
- When the resources become available then we will start up Matlab
- Commands to run:

```
module load slurm/summit  
sinteractive --reservation=rmacc
```
- Once we receive a prompt, then:

```
module load matlab  
matlab
```
- Once we finish we must exit!

Questions?

- Email rc-help@colorado.edu
- Link to survey on this topic:
<http://tinyurl.com/curc-survey16>

Speaker: Shelley Knuth

Title: How to Access Summit RMACC 2017

- Slides: <https://github.com/ResearchComputing/RMACC/2017>