



INTEL[®] THREADING BUILDING BLOCKS (INTEL[®] TBB) 2017

Multi-threading and heterogeneous computing made easy with Intel TBB


What is Intel® TBB?

Intel TBB is a highly templated C++ library designed to simplify the task of adding parallelism to your application by taking advantage of all the CPU's either on a single device or across multiple devices (heterogeneity).

Why should you use Intel® TBB?

- High Performance
- Easy to use API's
- Faster Time To Market
- Production Ready

Optimized for



Supports



Addresses



How to get Intel® TBB?

[Intel Parallel Studio XE](#)
[Intel System Studio](#)
[Free Tools Program](#)
[Open source site](#)

Applications

- Animation Rendering
- Numeric weather prediction
- Oceanography & Astrophysics
- Artificial Intelligence & Automation
- Genetic Engineering
- Medical applications (Image processing, MRI reconstruction)
- Remote sensing applications
- Socio Economics
- Financial sector (stock derivative pricing, statistics)
- Bulk updating data files
- Any Big Data problems

Find out more at: <http://software.intel.com/intel-tbb>

Contact us through our forum:

<http://software.intel.com/en-us/forums/intel-threading-building-blocks>

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.



Customer Success Stories

Dreamworks Fur Shader used Intel® TBB which produced an average of **5x speedup** on fur generation loop



Intel's TBB was an invaluable help in multi-threading our in-house renderer CGISudio and is now also used in animation and simulation software. Beside the **ease of use**, it takes care of the two most important aspects of running an application on multiple cores - **load balancing and scalability**.

- Maurice van Swaaij, Blue Sky Studios

"Intel® TBB provided us with **optimized code** that we did not have to develop or maintain for critical system services. I could assign my developers to code what we bring to the software table."


- Michaël Rouillé, CTO, Golaem

"Using Intel TBB's new flow graph feature, we accomplished what was previously not possible, **parallelize a very sizable task** graph with thousands of interrelationships – all in about a week.

- Robert Link
GCAM Project Scientist, Pacific Northwest National Laboratory


What Kind of applications needs to be multi-threaded?

Multi-threading is for applications where the problem can be broken down into tasks that can be run in parallel or the problem itself is massively parallel as some mathematics or analytical problems are.



Numeric Weather Prediction

Mathematical modelling
Forecasting the future state of weather
Data assimilation




Oceanography & Astrophysics

Study the wealth of ocean
PIC, PM and n-body simulations
Astrophysics research



Socio Economics

Modelling economy of nation/world
Scenario calculations & optimizations of economic models



Finite Element Analysis

Multi-physics problems
Design of huge structures like ships, dams, supersonic jets.
Solving Partial Differential Equations (PDE"S)




AI & Automation

Image Processing
Expert Systems
Natural Language Processing(NLP)
Pattern Recognition
ADAS



Genetic Engineering

DNA Sequence Analysis




Seismic Exploration

Use seismic waves to estimate earth's properties
Sensor Analysis



Weapon Research & Defense


Nuclear weapon's performance & certification
Plutonium research



Remote Sensing Applications

Used in Agriculture/Forestry to read specialized formats containing:

- sensor image data
- Georeferencing information
- Sensor metadata



Medical Applications

Medical Image Processing
Scanning human body/brain
MRI reconstruction
Vertebra detection & segmentation in X-ray images
Brain Fiber tracking



Energy Resource Exploration

Gather and manage information about energy resources like oil, natural gas
Maintain records of global energy crisis, nuclear reactor safety

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Advantages of using Intel TBB over other threading models




- Specify tasks instead of manipulating threads. Intel® TBB maps your logical tasks onto threads with full support for nested parallelism
- Intel TBB uses proven , efficient parallel patterns.
- Intel TBB uses work stealing to support the load balance of unknown execution time for tasks. This has the advantage of low-overhead polymorphism.
- Flow graph feature in Intel TBB allows developers to easily express dependency and data flow graphs.
- Has high level parallel algorithms and concurrent containers and low level building blocks like scalable memory allocator , locks and atomic operations.

Intel TBB – “ The backbone of CAD Exchanger’s parallelism”

Parallelism Brings CAD Exchanger* Software Dramatic Gains in Performance and User Satisfaction, Plus a Competitive Advantage “CAD Exchanger* is broadly using multi-threaded algorithms to increase performance on multi-core systems,” said Roman Lygin of CADEX, Ltd. “This is the key advantage over our competitors.” Benchmarks show how it outperforms earlier editions in significant ways:

- Some heavyweight computational algorithms, such as blended surface approximation, were **accelerated by 15X** over single-thread mode.
- Multi-threaded visualization significantly increased the responsiveness of the GUI application, which in turn improved the user experience. **Less time spent waiting** means more time to interact and innovate.
- Parallel file I/O is **2.5× faster**, and visualization time was **reduced by up to 4X**




Case Study link [here](#)

Intel TBB helps John Hopkins University prepare for a many core future

Modern DNA sequencing provides an inexpensive and high-resolution window into diverse aspects of biology, genetics, and disease. Like a microscope, a sequencer produces a snapshot of a collection of cells. Unlike a microscope, a sequencer does not provide a finished, ready-to-interpret image. Rather, it produces billions of tiny snippets (reads) of DNA that must first be composed into longer, interpretable units such as genes or chromosomes. Bowtie* and Bowtie 2* are widely used software tools produced in the University's Langmead Lab that allow biologists to piece together the fragmentary evidence generated by DNA sequencers.

Johns Hopkins and Intel have been collaborating on the Bowtie 2 application. Adding parallelism via Intel TBB resulted in a substantial speedup of the application. By splitting reads from parsing in a critical section, the team saw essentially ideal **scaling up to 120 threads**.

The team was able to **effectively prepare these core genomics software tools for the many-core future around the corner**.



Case Study link [here](#)

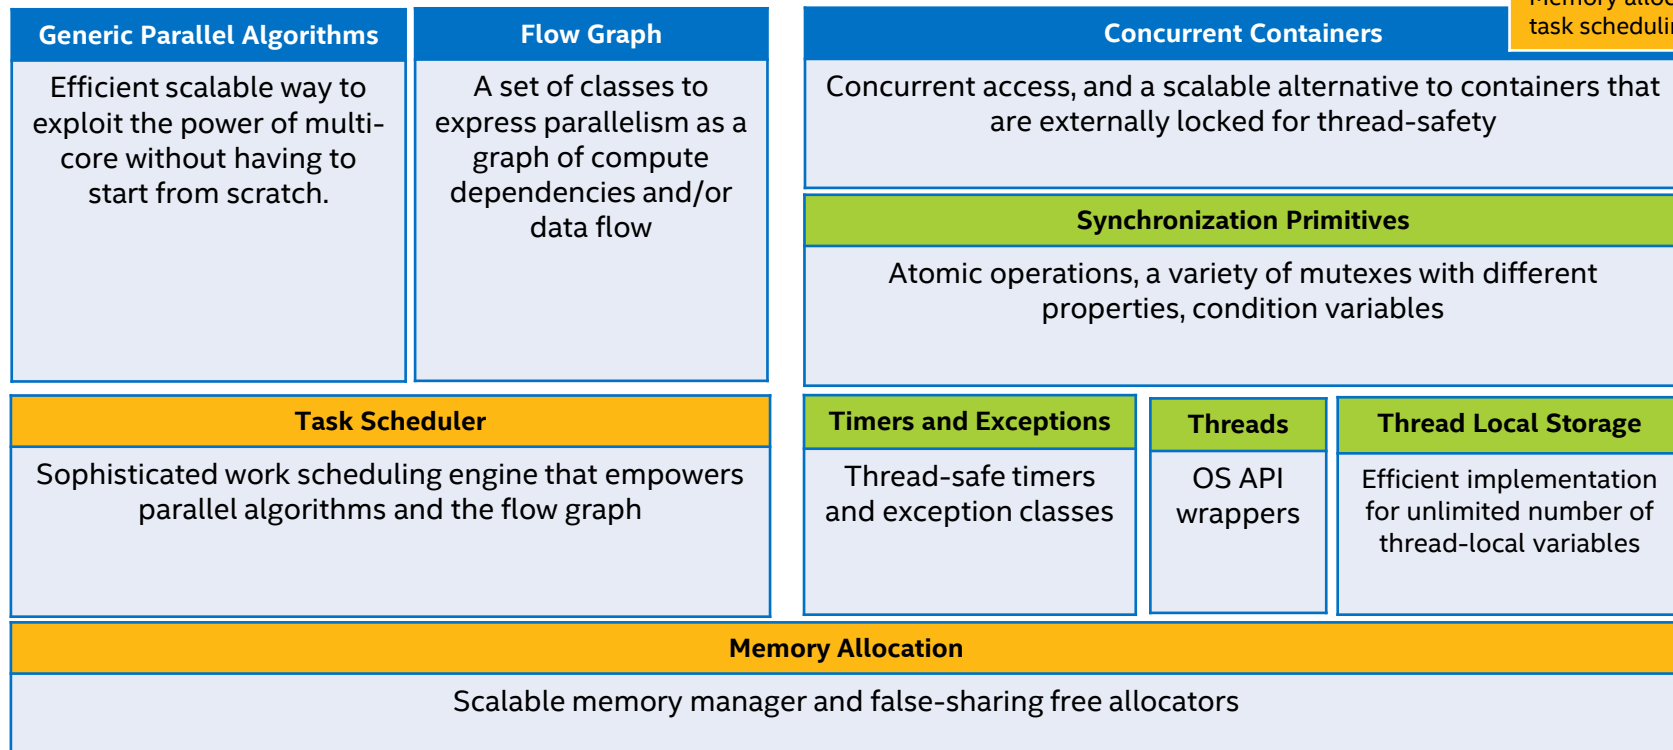
Rich Feature Set for Parallelism

Intel® Threading Building Blocks (Intel® TBB)

Parallel algorithms and
data structures

Threads and
synchronization

Memory allocation and
task scheduling



Features and Functions List

Intel® Threading Building Blocks (Intel® TBB)

Parallel algorithms and data structures

Threads and synchronization

Memory allocation and task scheduling

Generic Parallel Algorithms

- parallel_for
- parallel_reduce
- parallel_for_each
- parallel_do
- parallel_invoke
- parallel_sort
- parallel_deterministic_reduce
- parallel_scan
- parallel_pipeline
- pipeline

Flow Graph

- graph
- continue_node
- source_node
- function_node
- multifunction_node
- overwrite_node
- write_once_node
- limiter_node
- buffer_node
- queue_node
- priority_queue_node
- sequencer_node
- broadcast_node
- join_node
- split_node
- indexer_node

Concurrent Containers

- concurrent_unordered_map
- concurrent_unordered_multimap
- concurrent_unordered_set
- concurrent_unordered_multiset
- concurrent_hash_map
- concurrent_queue
- concurrent_bounded_queue
- concurrent_priority_queue
- concurrent_vector
- concurrent_lru_cache (preview)

Synchronization Primitives

- atomic
- mutex
- recursive_mutex
- spin_mutex
- spin_rw_mutex
- speculative_spin_mutex
- speculative_spin_rw_mutex
- queuing_mutex
- queuing_rw_mutex
- null_mutex
- null_rw_mutex
- reader_writer_lock
- critical_section
- condition_variable
- aggregator (preview)

Task Scheduler

- task
- task_group
- structured_task_group
- task_group_context
- task_scheduler_init
- task_scheduler_observer
- task_arena

Timers and Exceptions

- tick_count
- tbb_exception
- captured_exception
- movable_exception

Threads

- thread

Thread Local Storage

- combinable
- enumerable_thread_specific

Memory Allocation

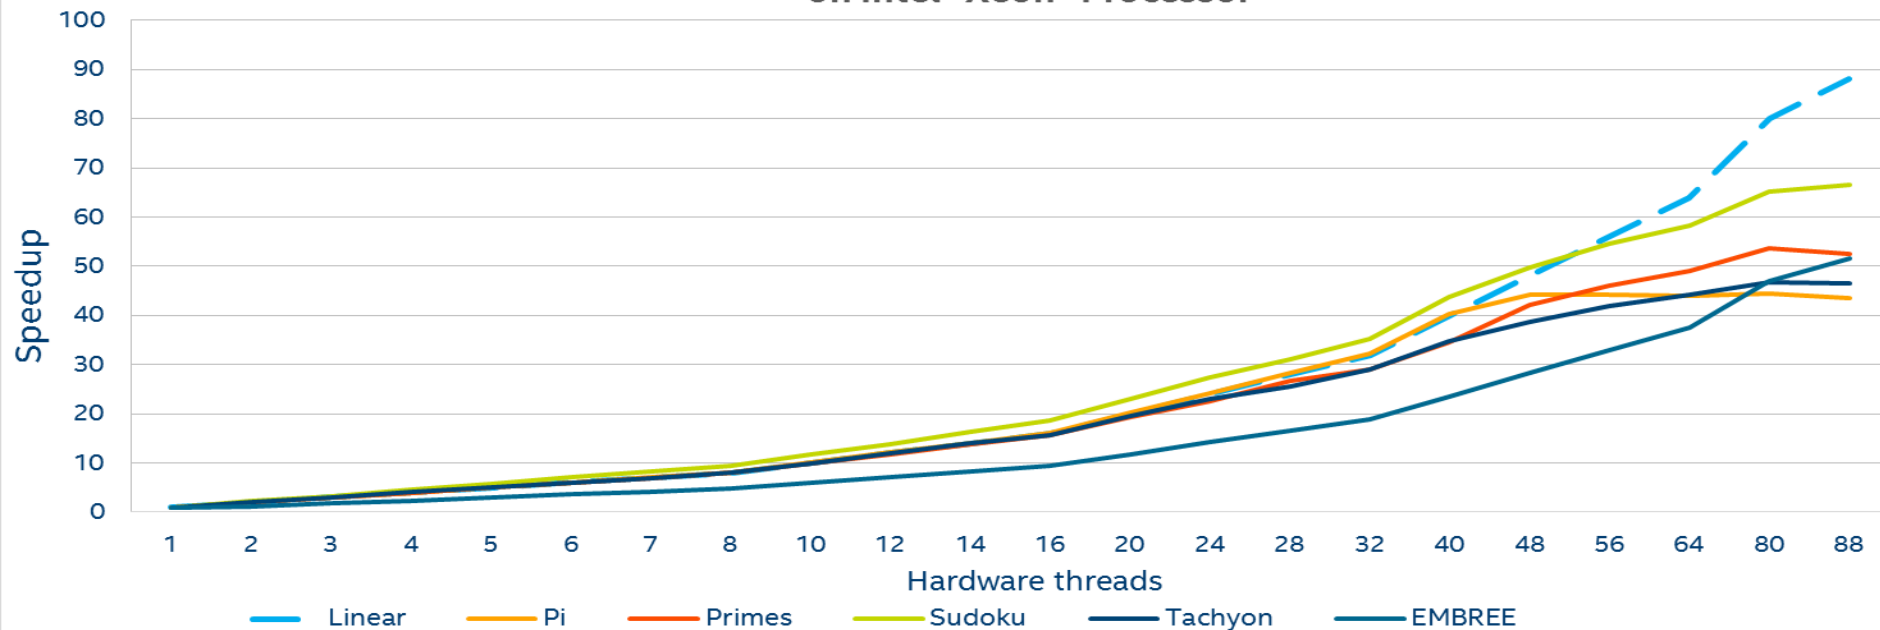
- tbb_allocator
- cache_aligned_allocator
- aligned_space
- scalable_allocator
- zero_allocator
- memory_pool (preview)

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.
*Other names and brands may be claimed as the property of others.



Excellent Performance Scalability with Intel® Threading Building Blocks 2017 on Intel® Xeon® Processor

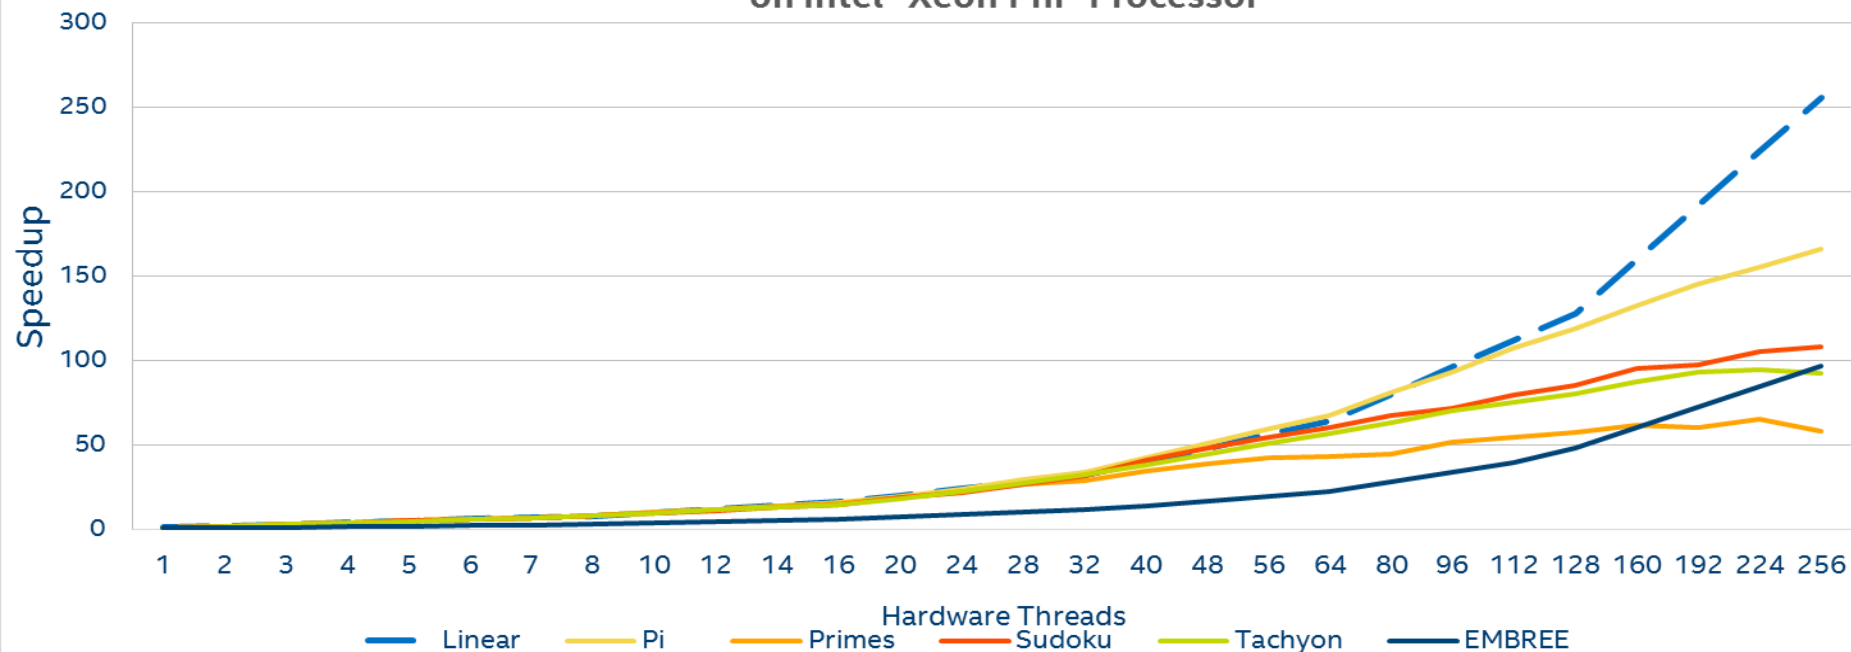


Configuration Info: Software Versions: Intel® C++ Intel® 64 Compiler, Version 17.0, Intel® Threading Building Blocks (Intel® TBB) 2017; Hardware: Intel(R) Xeon(R) CPU E5-2699 v4 @ 2.20GHz 44/88T), 128GB Main Memory; Operating System: Red Hat Enterprise Linux Server release 7.2 (Maipo), kernel 3.10.0-327.4.5.el7.x86_64; Benchmark Source: Intel Corp. Note: sudoku, primes and tachyon are included with Intel TBB. Performance tests and ratings are measured using specific computer systems and/or components and reflect the approximate performance of Intel products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance. Buyers should consult other sources of information to evaluate the performance of systems or components they are considering purchasing. For more information on performance tests and on the performance of Intel products, refer to www.intel.com/performance/resources/benchmark_limitations.htm.

* Other brands and names are the property of their respective owners.

Optimization Notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice. Notice revision #20110804

Excellent Performance Scalability with Intel® Threading Building Blocks 2017 on Intel® Xeon Phi® Processor



Configuration Info: Software Versions: Intel® C++ Intel® 64 Compiler, Version 17.0, Intel® Threading Building Blocks (Intel® TBB) 2017; Hardware: KNL (Intel(R) Xeon Phi(TM) CPU 7250 @ 1.40GHz (68C/272T)), 128GB Main Memory; Operating System: Red Hat Enterprise Linux Server release 7.2 (Maipo), kernel 3.10.0-327.13.1.el7.mpsp_1.3.2.100.x86_64; Benchmark Source: Intel Corp. Note: sudoku, primes and tachyon are included with Intel TBB. Performance tests and ratings are measured using specific computer systems and/or components and reflect the approximate performance of Intel products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance. Buyers should consult other sources of information to evaluate the performance of systems or components they are considering purchasing. For more information on performance tests and on the performance of Intel products, refer to www.intel.com/performance/resources/benchmark_limitations.htm.

* Other brands and names are the property of their respective owners.

Optimization Notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice. Notice revision #20110804


Optimization Notice

Copyright © 2015, Intel Corporation. All rights reserved.

*Other names and brands may be claimed as the property of others.



Task Execution in Intel TBB




(A simplified version of the scheduler)

Generic algorithms allow reuse of proven parallel patterns

Intel® Threading Building Blocks (Intel® TBB)

Sequential version

```
int mandel(Complex c, int max_count) {  
    int count = 0; Complex z = 0;  
    for (int i = 0; i < max_count; i++) {  
        if (abs(z) >= 2.0) break;  
        z = z*z + c; count++;  
    }  
    return count;  
}
```



```
for (int i = 0; i < max_row; i++) {  
    for (int j = 0; j < max_col; j++) {  
        p[i][j] = mandel( Complex(scale(i), scale(j)), depth);  
    }  
}
```

For each point, is $z = z^2 + c$ bounded?

Mandelbrot Speedup


Intel® Threading Building Blocks (Intel® TBB)

```
int mandel(Complex c, int max_count) {  
    int count = 0; Complex z = 0;  
    for (int i = 0; i < max_count; i++) {  
        if (abs(z) >= 2.0) break;  
        z = z*z + c; count++;  
    }  
    return count;  
}
```

Parallel algorithm


```
parallel_for( 0, max_row,  
    [&](int i) {  
        for (int j = 0; j < max_col; j++)  
            p[i][j]=mandel(Complex(scale(i),scale(j)),depth);  
    }  
);
```

Use C++ lambda functions to define function object in-line




Task is a function object


A parallel_for recursively divides the range into subranges that execute as tasks - Intel® Threading Building Blocks (Intel® TBB)



Split range...




.. recursively...



...until \leq
grainsize.

A parallel_for recursively divides the range into subranges that execute as tasks - Intel® Threading Building Blocks (Intel® TBB)




Split range...



.. recursively...




...until \leq
grainsize.



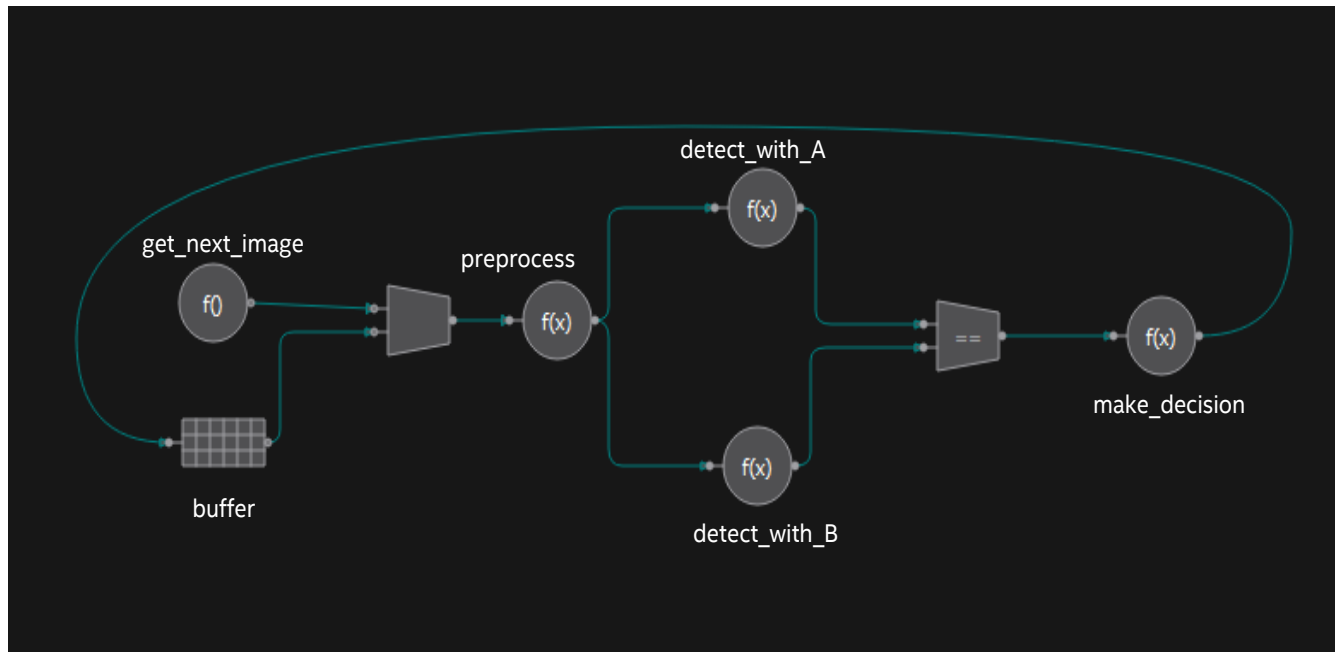
Flow Graph Hello World Example

Users create nodes and edges, interact with the graph and wait for it to complete

```
tbb::flow::graph g;  
tbb::flow::continue_node< tbb::flow::continue_msg >  
    h( g, []( const continue_msg & ) { std::cout << "Hello "; } );  
tbb::flow::continue_node< tbb::flow::continue_msg >  
    w( g, []( const continue_msg & ) { std::cout << "World\n"; } );  
tbb::flow::make_edge( h, w );  
h.try_put(continue_msg());  
g.wait_for_all();
```

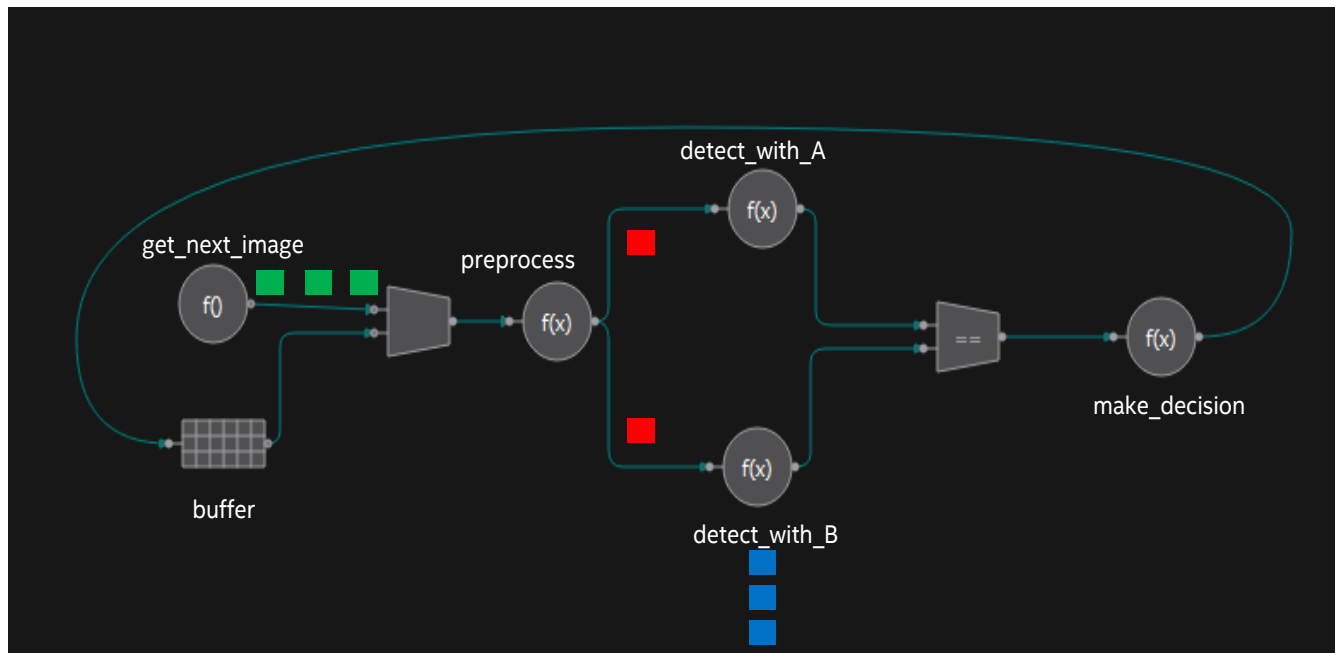


An example feature detection algorithm



Can express **pipelining**, **task parallelism** and **data parallelism**


An example feature detection algorithm



Can express **pipelining**, **task parallelism** and **data parallelism**

And supports nested parallelism with Intel TBB, OpenMP,
Intel® Cilk™ Plus, Intel® Math Kernel Library (Intel® MKL), etc...*

CPU Programming Model Hierarchy



- **Message Driven (TBB Flow Graph)**
Uses same resources/scheduler as (B) since (A) is just another hierarchical layer
- **Fork Join / Tasking (TBB Tasks)**
Tolerant of unanticipated CPU loads and support efficient composition
- **SIMD**
Requires compiler support. New standardization proposal for parallel STL in C++ will integrate this layer into the same software stack.

Intel® Threading Building Blocks (Intel® TBB) is the C++ library that provides what is needed for the **Message Driven** and **Fork Join / Tasking** layers

Use all your available compute resources across HW and SW through Intel TBB

Hardware


Integrated graphics, media, CPU's along with discrete co-processors & accelerators (FPGA's, fixed function devices etc)



+

Software

Other threading as well as domain specific libraries and API's



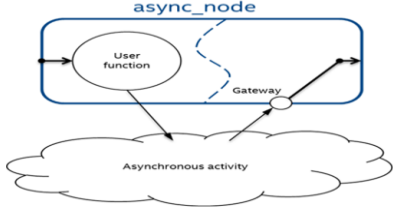
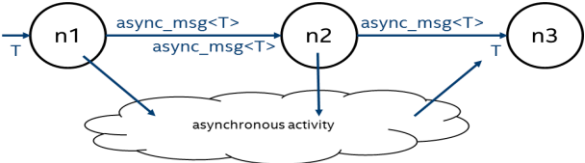
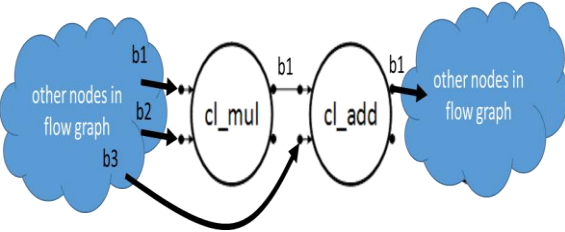
Composability layer with Intel TBB

One threading engine under all hardware (CPU) side work

Co-ordination layer with Intel TBB flow graph

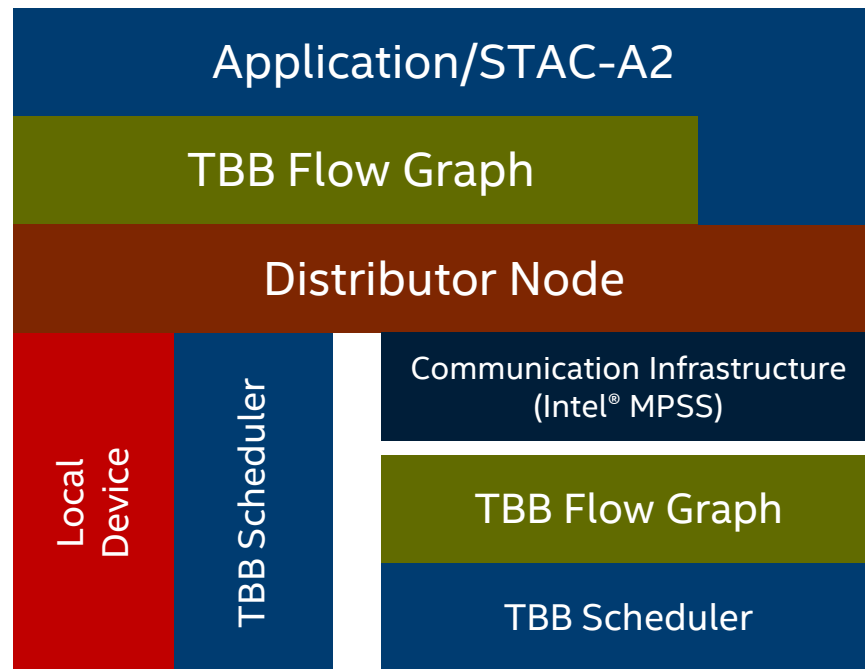
Be the glue connecting HW & SW, expose parallelism between blocks & simplify integration

Heterogeneous Features supported today through Intel TBB flow graph API

Feature	Description	Diagram
<i>Async_node</i>	Allows flow graph to offload data to any asynchronous activity and receive it back to continue execution on CPU.	 A diagram showing an 'async_node' box. Inside the box, an arrow points from the left into an oval labeled 'User function'. Another arrow points from the 'User function' to a small circle labeled 'Gateway'. A dashed arrow points from the 'Gateway' back to the 'User function'. An arrow points from the 'Gateway' out of the box to a cloud labeled 'Asynchronous activity'. Another arrow points from the cloud back to the 'Gateway'.
<i>Async_msg</i>	Enables async communication with chaining across graph nodes. User manages communication. Graph runs on host.	 A diagram showing three nodes in a circle, labeled n1, n2, and n3. An arrow labeled 'T' enters n1 from the left. An arrow labeled 'async_msg<T>' points from n1 to n2. Another arrow labeled 'async_msg<T>' points from n2 to n3. An arrow labeled 'T' exits n3 to the right. Below the nodes is a cloud labeled 'asynchronous activity'. Arrows point from n1, n2, and n3 down to the cloud. An arrow points from the cloud up to n3.
<i>Opencl_node</i>	A specialization of heterogeneous node for openCL. User provide openCL program & Kernel, runtimes handles initialization, buffer management, comms, Graph runs on host	 A diagram showing two clouds on the left and right, both labeled 'other nodes in flow graph'. Between them are two ovals labeled 'cl_mul' and 'cl_add'. An arrow labeled 'b1' points from the left cloud to 'cl_mul'. Two arrows labeled 'b2' and 'b3' point from the left cloud to 'cl_mul'. An arrow labeled 'b1' points from 'cl_mul' to 'cl_add'. An arrow labeled 'b1' points from 'cl_add' to the right cloud. A curved arrow points from the bottom of the left cloud to the bottom of 'cl_add'.

STAC-A2 Implementation of Intel TBB's Heterogeneous feature

- Compute-intensive analytic workload involved in pricing and risk management
- Implemented with Intel® TBB flow graph, TBB parallel algorithms and OpenMP vectorization
- Uses asynchronous support in flow graph, distributor_node, to offload to the Intel® Xeon Phi coprocessor
- Using a token-based system enables dynamic load balancing between CPU and coprocessor
- Which Greeks are calculated on which resource is determined dynamically based on resource availability



<https://stacresearch.com/news/2015/11/03/stac-report-stac-a2-system-dual-xeon-phi-cards>

Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2016, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

