

INTEL® VTUNE™ AMPLIFIER PERFORMANCE PROFILER

Faster, Scalable Code, Faster

Intel® VTune™ Amplifier Performance Profiler

Accurate Data - Low Overhead

CPU, GPU, FPU, threading, bandwidth...

Meaningful Analysis

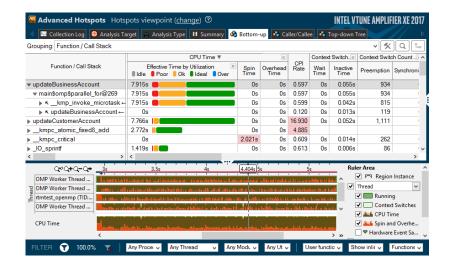
- Threading, OpenMP region efficiency
- Memory access, storage device

Easy

- Data displayed on the source code
- Easy set-up, no special compiles

"Last week, Intel® VTune™ Amplifier helped us find almost 3X performance improvement. This week it helped us improve the performance another 3X."

Claire Cates
Principal Developer
SAS Institute Inc.



http://intel.ly/vtune-amplifier-xe



Two Great Ways to Collect Data

Intel[®] VTune[™] Amplifier

Software Collector	Hardware Collector		
Uses OS interrupts	Uses the on chip Performance Monitoring Unit (PMU)		
Collects from a single process tree	Collect system wide or from a single process tree.		
~10ms default resolution	~1ms default resolution (finer granularity - finds small functions)		
Either an Intel® or a compatible processor	Requires a genuine Intel® processor for collection		
Call stacks show calling sequence	Optionally collect call stacks		
Works in virtual environments	Works in a VM only when supported by the VM		
	(e.g., vSphere*, KVM)		
No driver required	Requires a driver - Easy to install on Windows - Linux requires root (or use default perf driver)		

No special recompiles - C, C++, C#, Fortran, Java, Assembly



A Rich Set of Performance Data

Intel® VTune™ Amplifier

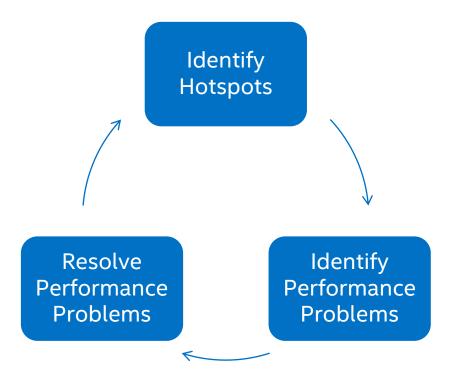
Software Collector	Hardware Collector		
Basic Hotspots Which functions use the most time?	Advanced Hotspots Which functions use the most time? Where to inline? – Statistical call counts		
Concurrency Tune parallelism. Colors show number of cores used.	General Exploration Where is the biggest opportunity? Cache misses? Branch mispredictions?		
Locks and Waits Tune the #1 cause of slow threaded performance: – waiting with idle cores.	Advanced Analysis Memory-access, HPC Characterization, etc		
Any IA86 processor, any VM, no driver	Higher res., lower overhead, system wide		

No special recompiles - C, C++, C#, Fortran, Java, Assembly



PERFORMANCE ANALYSIS WITH INTEL® VTUNE™ AMPLIFIER

Analysis Workflow



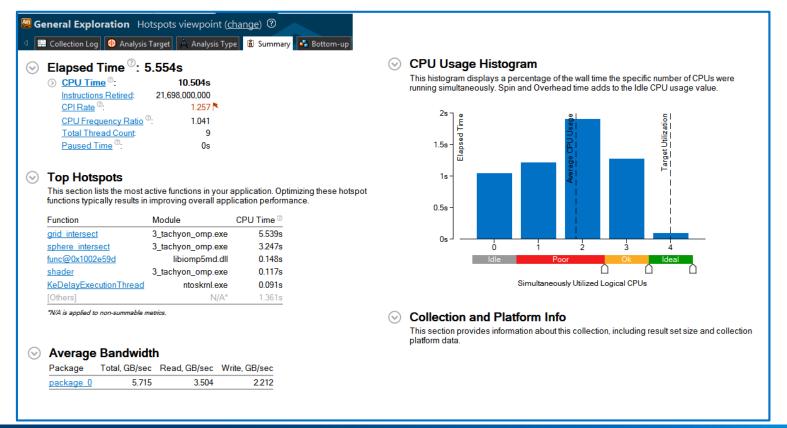


VTUNE BY EXAMPLE

VTUNE FEATURES

Example: Hotspots Analysis

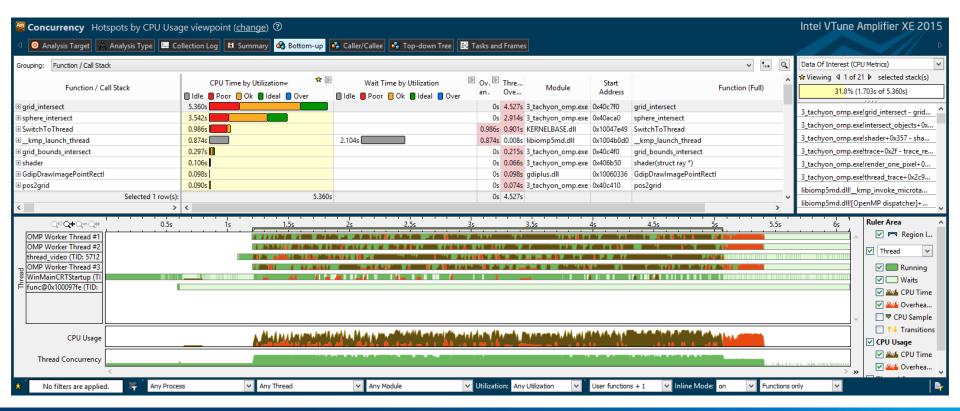
Summary View





Example: Concurrency Analysis

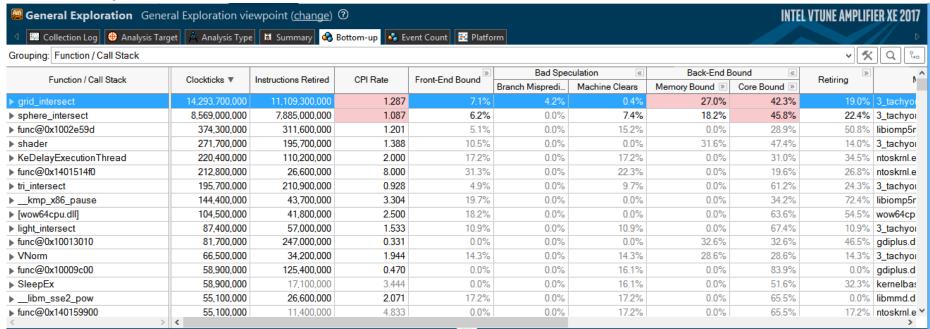
Bottom-up View





Example: General Exploration

Bottom-up View



Find real CPU stalls due to cache misses, instruction fetch misses, branch misprediction, and a lot more.

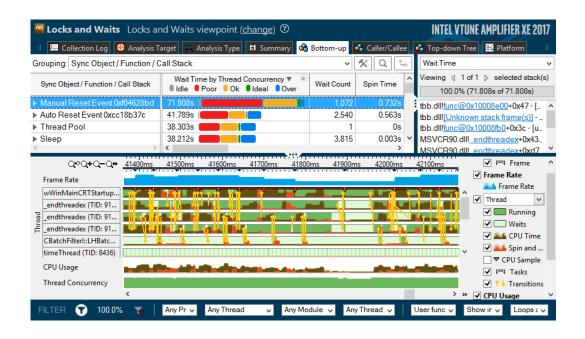


Intel[®] VTune[™] Amplifier

Tune Applications for Scalable Multicore Performance

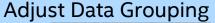
Agenda

- Data Collection –
 Rich set of performance data
- Data Analysis -Find answers fast
- Flexible workflow
 - User i/f and command line
 - Compare results
 - Remote collection
- Additional Features
- Summary



Find Answers Fast

Intel[®] VTune[™] Amplifier



Function - Call Stack

Module - Function - Call Stack

Source File - Function - Call Stack

Thread - Function - Call Stack

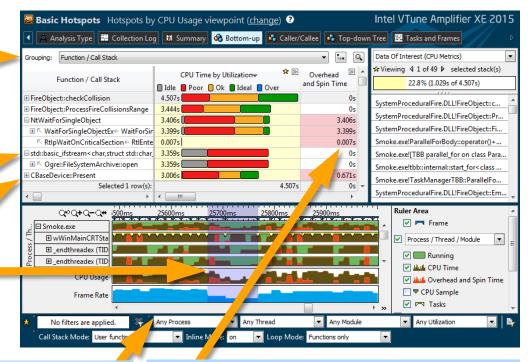
... (Partial list shown)

Double Click Function to View Source

Click [+] for Call Stack

Filter by Timeline Selection (or by Grid Selection)





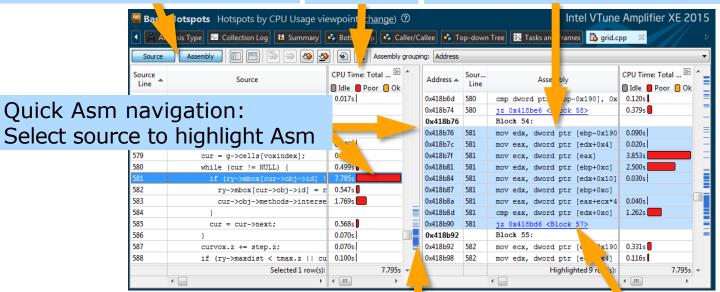
Filter by Process & Other Controls

Tuning Opportunities Shown in Pink. Hover for Tips

See Profile Data On Source / Asm

Double Click from Grid or Timeline

View Source / Asm or both CPU Time Right click for instruction reference manual



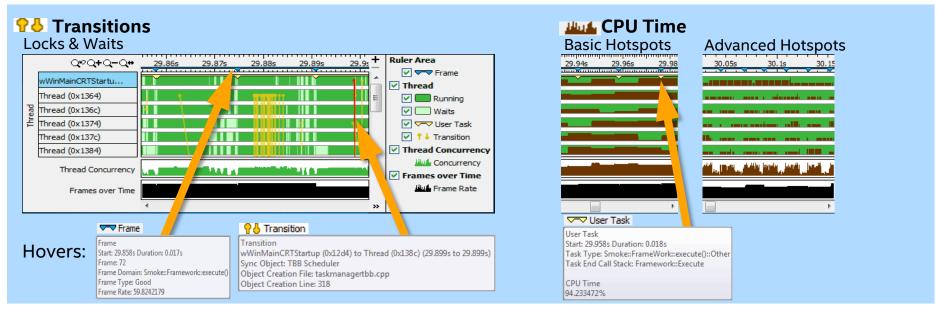
Scroll Bar "Heat Map" is an overview of hot spots

Click jump to scroll Asm



Timeline Visualizes Thread Behavior

Intel® VTune™ Amplifier



Optional: Use API to mark frames and user tasks Frame Suser Task



Optional: Add a mark during collection



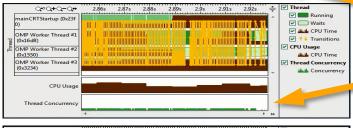
Visualize Parallel Performance Issues

Look for Common Patterns

Coarse Grain Locks CPU Usage

Thread Concurrency

High Lock Contention



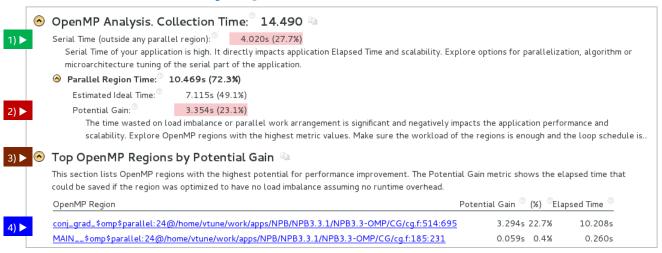
Low Concurrency

Load Imbalance



Tune OpenMP for Efficiency and Scalability

Fast Answers: Is My OpenMP Scalable? How Much Faster Could It Be?



The summary view shown above gives fast answers to four important OpenMP tuning questions:

- 1) Is the serial time of my application significant enough to prevent scaling?
- 2) How much performance can be gained by tuning OpenMP?
- 3) Which OpenMP regions / loops / barriers will benefit most from tuning?
- 4) What are the inefficiencies with each region? (click the link to see details)



Intel[®] VTune[™] Amplifier

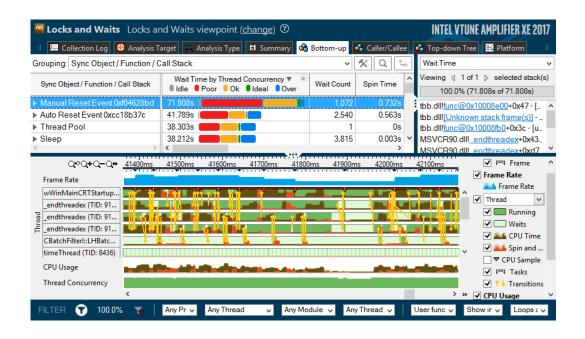
Tune Applications for Scalable Multicore Performance

Agenda

- Data Collection –
 Rich set of performance data
- Data Analysis Find answers fast



- User i/f and command line
- Compare results
- Remote collection
- Additional Features
- Summary



Command Line Interface

Automate analysis

amplxe-cl is the command line:

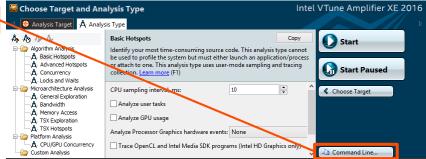
- -Windows: C:\Program Files (x86)\Intel\VTune Amplifier XE \bin[32|64]\amplxe-cl.exe
- -Linux: /opt/intel/vtune_amplifier_xe/bin[32|64]/amplxe-cl

Help: amplxe-cl -help



Use UI to setup

- 1) Configure analysis in UI
- 2) Press "Command Line..." button
- 3) Copy & paste command



Great for regression analysis – send results file to developer Command line results can also be opened in the UI

MPI Analysis

Command line:

```
> mpirun -n 16 -ppn 4 -l amplxe-cl -collect advanced-hotspots - trace-mpi -result-dir my_result -- my_app.a
```

Or use gtool:

```
> mpirun -gtool "amplxe-cl -collect memory-access -result-dir
my_result:7,5" my_app.a
```

Each process data is presented for each node they were running on:

```
my_result.host_name1 (rank 0-3)
my_result.host_name2 (rank 4-7)
my_result.host_name3 (rank 8-11)
my_result.host_name4 (rank 12-15)
```



Interactive Remote Data Collection

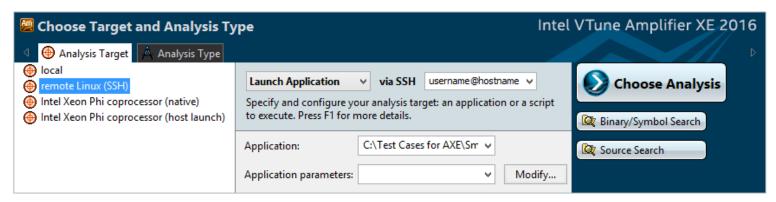
Performance analysis of remote systems just got a lot easier

Interactive analysis

- 1) Configure SSH to a remote Linux* target
- 2) Choose and run analysis with the UI

Command line analysis

- Run command line remotely on Windows* or Linux* target
- 2) Copy results back to host and open in UI



Conveniently use your local UI to analyze remote systems



Compare Results Quickly - Sort By Difference

Intel[®] VTune[™] Amplifier

Quickly identify cause of regressions.

- Run a command line analysis daily
- Identify the function responsible so you know who to alert

Compare 2 optimizations – What improved?

Compare 2 systems – What didn't speed up as much?

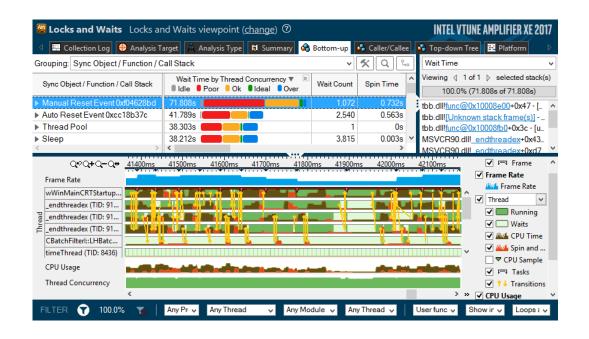
Grouping: Function / Call Stack ▼							
Function / Call Stack	CPU Time:Difference	Module	CPU Time:r	007hs 🌣	CPU Tim	ne:r006hs	<u>^</u>
■ FireObject::checkCollision	4.850s	SystemProceduralFire.DLL	6.281s		1.431s	0	
	4.644s	SystemProceduralFire.DLL	5.643s		0.999s	0	
	3.765s	RenderSystem_Direct3D9.DLL	9.184s		5.419s		

Intel[®] VTune[™] Amplifier

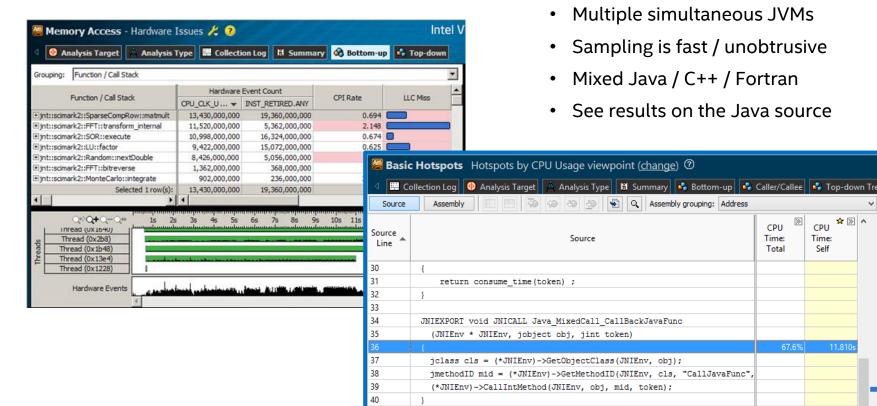
Tune Applications for Scalable Multicore Performance

Agenda

- Data Collection –
 Rich set of performance data
- Data Analysis Find answers fast
- Flexible workflow
 - User i/f and command line
 - Compare results
 - Remote collection
- Additional Features
 - Summary



Java Analysis





Optimize Private Cloud-Based Applications

Profile Enterprise Applications

- ■Native C, C++, Fortran*
- Attach to running Java* services (e.g., Mail)
- Profile Java daemons without restart

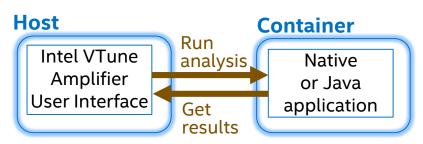
Accurate, Low-Overhead Data Collection

- Advanced hotspots and hardware events
- •Memory analysis
- •Accurate stack information for Java and HHVM*

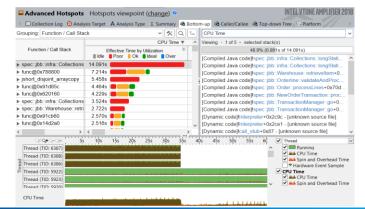
Popular Containers Supported

- Docker*
- ■Mesos*

Software collectors (e.g., locks & waits) and Python* profiling are not currently available for containers.



- No container configuration required
- Detection of the container is automatic



Optimize Memory Access

Memory Access Analysis - Intel® VTune™ Amplifier 2017

Tune data structures for performance

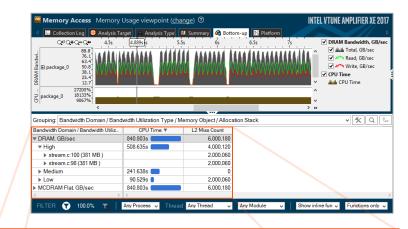
- Attribute cache misses to data structures (not just the code causing the miss)
- Support for custom memory allocators

Optimize NUMA latency & scalability

- True & false sharing optimization
- Auto detect max system bandwidth
- Easier tuning of inter-socket bandwidth

Easier install, Latest processors

- No special drivers required on Linux*
- Intel® Xeon Phi™ processor MCDRAM (high bandwidth memory) analysis



Bandwidth Domain / Bandwidth Utiliz	CPU Time ▼	L2 Miss Count
▼ DRAM, GB/sec	840.803s	6,000,180
▼ High	508.635s	4,000,120
▶ stream.c:100 (381 MB)		2,000,060
▶ stream.c:98 (381 MB)		2,000,060
▶ Medium	241.638s	0
▶ Low	90.529s	2,000,060
▶ MCDRAM Flat, GB/sec	840.803s	6,000,180

Storage Device Analysis (HDD, SATA or NVMe SSD)

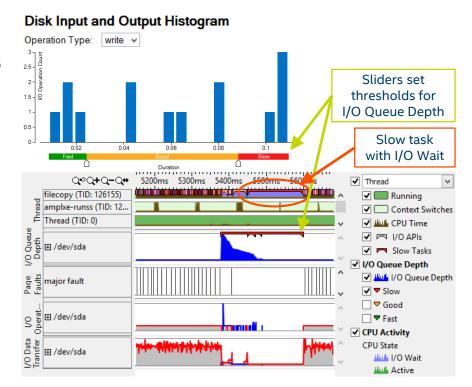
Intel® VTune™ Amplifier

Are You I/O Bound or CPU Bound?

- Explore imbalance between I/O operations (async & sync) and compute
- Storage accesses mapped to the source code
- See when CPU is waiting for I/O
- Measure bus bandwidth to storage

Latency analysis

- Tune storage accesses with latency histogram
- Distribution of I/O over multiple devices



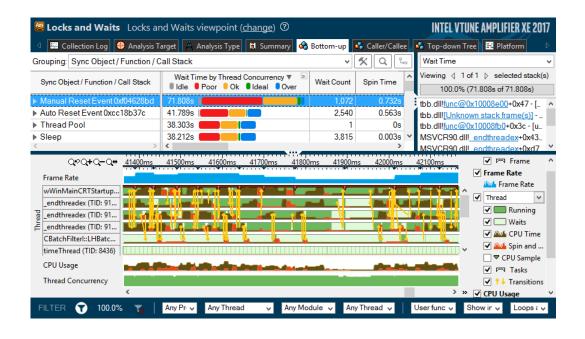
Intel[®] VTune[™] Amplifier

Tune Applications for Scalable Multicore Performance

Agenda

- Data Collection –
 Rich set of performance data
- Data Analysis Find answers fast
- Flexible workflow
 - User i/f and command line
 - Compare results
 - Remote collection
- Additional Features





Intel® VTune™ Amplifier

Faster, Scalable Code Faster

Get the Data You Need

- Hotspot (Statistical call tree), Call counts (Statistical)
- Thread Profiling Concurrency and Lock & Waits Analysis
- Cache miss, Bandwidth analysis...¹
- GPU Offload and OpenCL™ Kernel Tracing

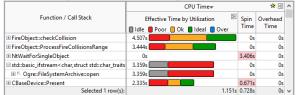
Find Answers Fast

- View Results on the Source / Assembly
- OpenMP Scalability Analysis, Graphical Frame Analysis
- Filter Out Extraneous Data Organize Data with Viewpoints
- Visualize Thread & Task Activity on the Timeline

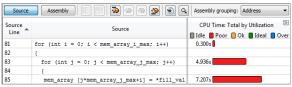
Easy to Use

- No Special Compiles C, C++, C#, Fortran, Java, ASM
- Visual Studio* Integration or Stand Alone
- Local & Remote Data Collection, Command Line
- Analyze Windows* & Linux* data on OS X*2

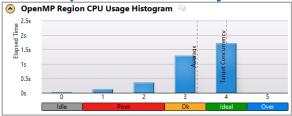
Quickly Find Tuning Opportunities



See Results On The Source Code



Tune OpenMP Scalability



Visualize & Filter Data





Legal Disclaimer & Optimization Notice

INFORMATION IN THIS DOCUMENT IS PROVIDED "AS IS". NO LICENSE, EXPRESS OR IMPLIED, BY ESTOPPEL OR OTHERWISE, TO ANY INTELLECTUAL PROPERTY RIGHTS IS GRANTED BY THIS DOCUMENT. INTEL ASSUMES NO LIABILITY WHATSOEVER AND INTEL DISCLAIMS ANY EXPRESS OR IMPLIED WARRANTY, RELATING TO THIS INFORMATION INCLUDING LIABILITY OR WARRANTIES RELATING TO FITNESS FOR A PARTICULAR PURPOSE, MERCHANTABILITY, OR INFRINGEMENT OF ANY PATENT, COPYRIGHT OR OTHER INTELLECTUAL PROPERTY RIGHT.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products.

Copyright © 2016, Intel Corporation. All rights reserved. Intel, Pentium, Xeon, Xeon Phi, Core, VTune, Cilk, and the Intel logo are trademarks of Intel Corporation in the U.S. and other countries.

Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804