

Future of HPC Software

Moderator: Michael Walker

Panelists:

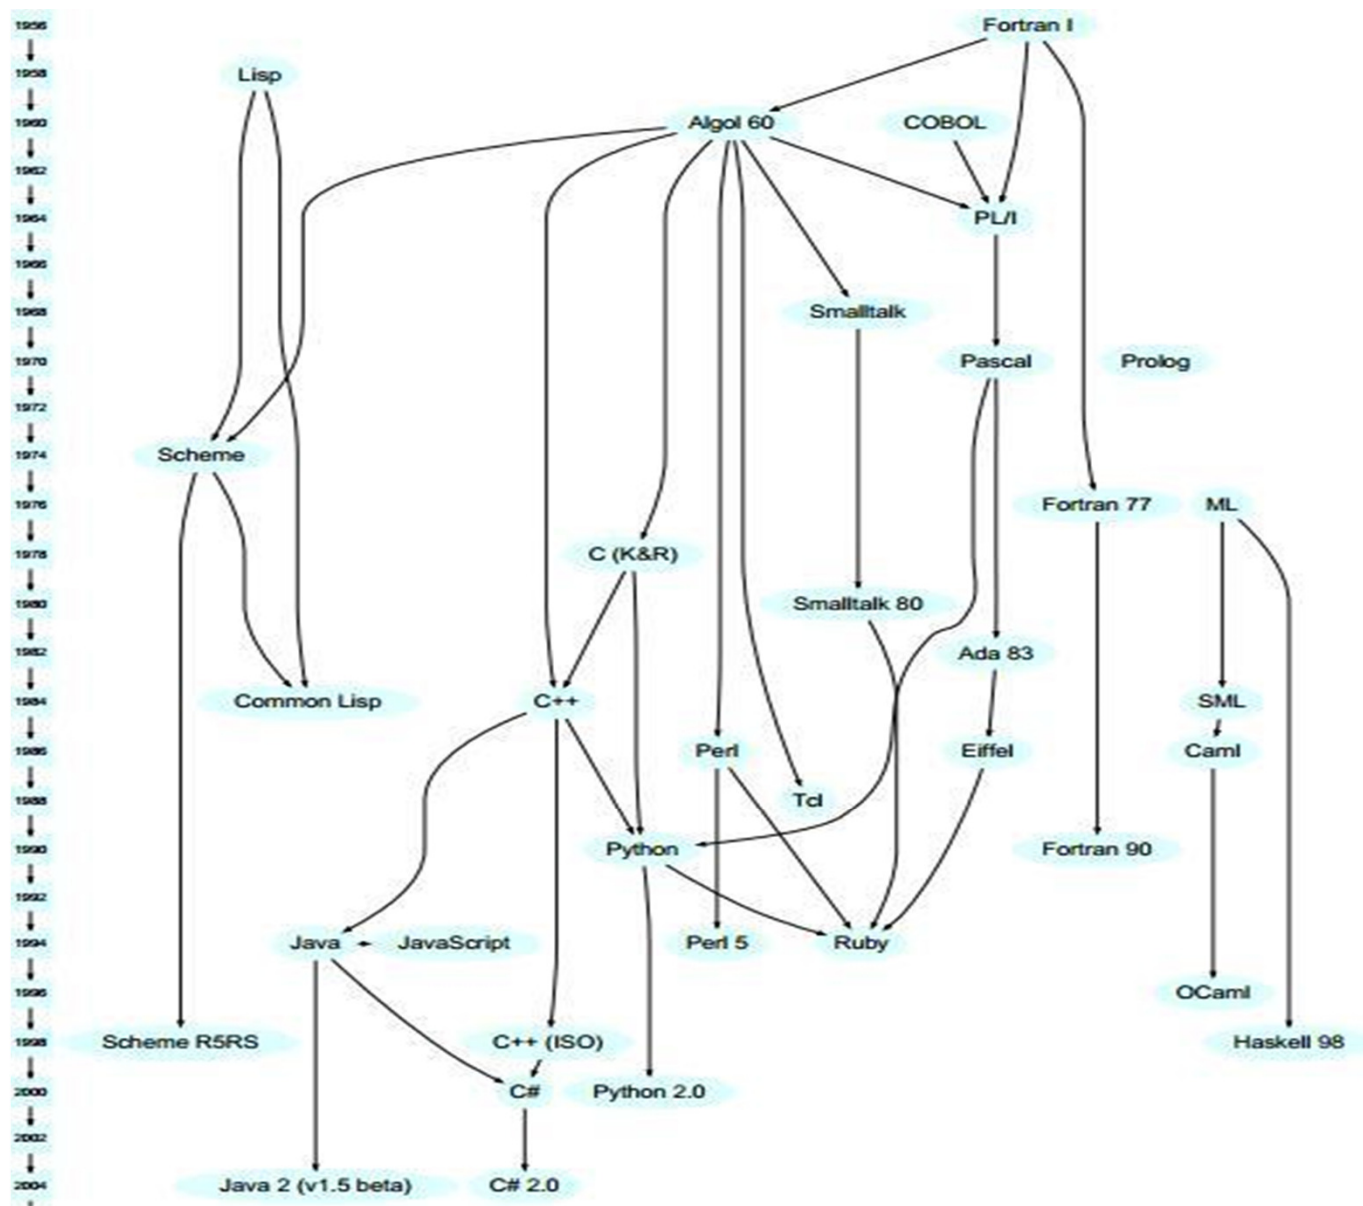
Monte Lunacek

Wesley Jones


Michael Malak

Cary Miller

Andrew Weekley



What programming/statistics languages you used for an analytics / data mining / data science work in 2013? [713 votes total]



Most popular languages:

R (61% of responders)

Python (39%)

SQL (37%)

Language with highest relative growth (2013 vs 2012)
was Julia, which doubled in popularity, but still was
used only by 0.7% in 2013.

Largest relative increases in share of usage from 2012 to 2013:

Pig Latin/Hive/other Hadoop-based languages, 19% growth, from 6.7% in 2012 to 8.0% in 2013

R, 16% growth

SQL, 14% growth (perhaps the result of increasing number of SQL interfaces to Hadoop and other Big Data systems?)

Languages with largest decline in share of usage:

Lisp/Clojure, 77% down

Perl, 50% down

Ruby, 41% down

C/C++, 35% down

Unix shell/awk/sed, 25% down

Java, 22% down






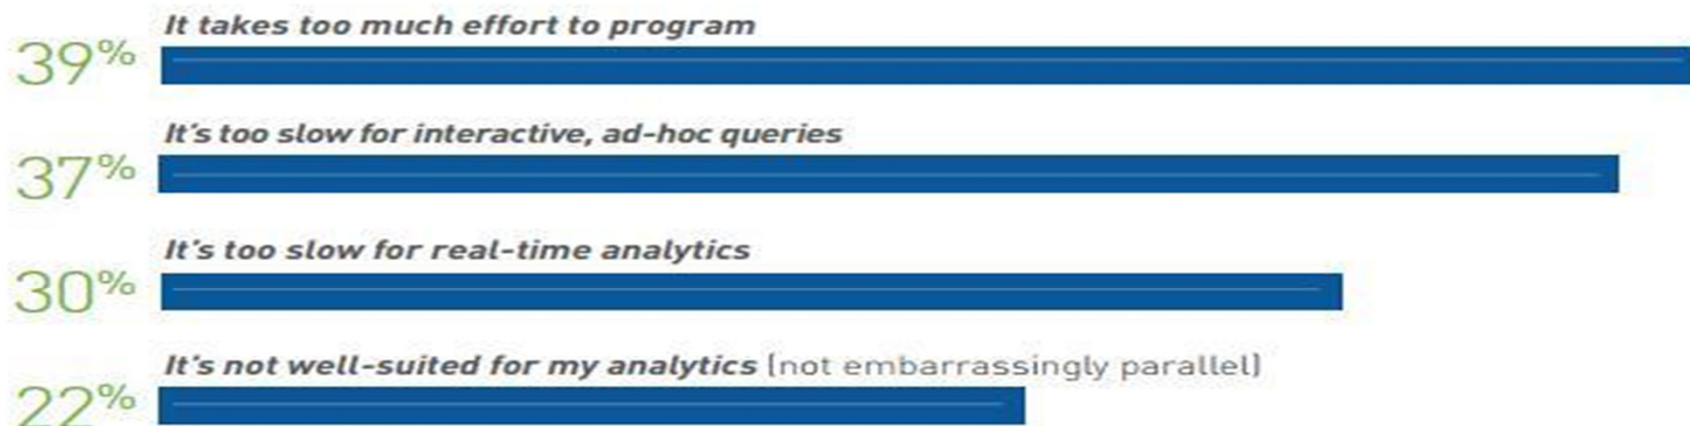
pythonTM



Top500 HPC Application Areas



From the 76% reporting problems, what are the limitations of Hadoop / SPARK?



of data scientists who tried Hadoop or
SPARK have stopped using it

New HPC language needed for performance at Petascale.

New HPC language needed for performance at Exascale.

Needed to improve software productivity.

New HPC language should:

- Address performance issues of large future systems
- Express well common HPC programming patterns
- Support well performance programming, with incremental code refinement
- Support OO
- Take advantage of advances in PL's: strong typing, type and memory safety, atomicity, efficient support for generic programming

New HPC language should:

- Take advantage of advances in compilers: dynamic compilation, heuristic search, telescoping languages
- Coexist with existing languages
- Provide state of the art ADE
- Largely built on language with large market

Non Technical Obstacles

- It takes money to make a good compiler; there is no market for HPC unique optimizations
- It takes time to make a good compiler; there is no funding mechanism for a sustained 5 years development effort
- It takes people to make a good compiler; there is no independent compiler company

Should Hardware vendors develop the HPC ADE?

Good ADE is more than language:

- Porting tools
- Good support for performance tuning
- Tools for refactoring
- Notation for capturing tuning decisions
- Good observability
- Integrated performance stream mining

Bottom Line

- HPC is hampered by lack of good software support
- Language is only part of the problem
- Most obstacles are not technological
- Key issues for petascale computing are just now being addressed